

Latent class analysis

Latent profile analysis a.k.a. Gaussian mixture modeling

DL Oberski & L Boeschoten

Learning goals



- Understand and apply basic latent profile analysis (a.k.a. Gaussian mixture modeling)
- Apply further clustering evaluation techniques
- Understand & apply component merging
- Use `mclust` in R

Multivariate model-based clustering


- With 2 observed features:
 - mean becomes a vector of 2 means
 - standard deviation turns into a 2x2 variance-covariance matrix determining the shape of the cluster
- So we have multiple within-cluster parameters:
 - Two means
 - Two variances, one for each observed variable
 - A single covariance among the features
- Together, the 11 parameters define the likelihood in bivariate space, which from the top looks like ellipses


Multivariate model-based clustering


$$p(\mathbf{y}|\boldsymbol{\theta}) = \pi_1^X MVN(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1) + (1 - \pi_1^X)MVN(\boldsymbol{\mu}_2, \boldsymbol{\Sigma}_2)$$




EM algorithm for Gaussian mixture model (LPA)









Multivariate model-based clustering

- Cluster shape parameters (the variance-covariance matrix) *can* be constrained to be *equal* across clusters
- Can also be *different* across clusters
- More flexible, complex model
- Think: **bias-variance tradeoff**

How to evaluate clustering results

1. Use of external information
2. Visual exploration
3. Stability assessment / sensitivity analysis
4. Internal validation indexes
- 5. Testing for clustering structure**

Much more info & helpful advice: Clustering strategy & method selection (ch 31 of Handbook of clustering), <https://arxiv.org/pdf/1503.02059.pdf>

File size increases with number clusters






Image loss decreases with number of clusters





- More clusters gives **better “fit”** in terms of reconstruction of the image (compression is less “lossy”)
- More clusters gives **bigger file size** (solution is more complex, takes more bytes to store)
- So the **model loss and model complexity trade off against each other**
- This is a common theme in (unsupervised) machine learning.

Model fit criteria

- BIC: “Schwarz/Bayesian information criterion”
- AIC: “Another/Akaike information criterion”
(*same as BIC but penalty is m*)
- AIC3: The same as AIC but penalty is $\frac{3}{2}m$
- ICL: “Integrated information criterion” (Biernacki et al. 2000)
(*Same as BIC but penalized by entropy of classification*)
- (Others based on):
 - *Minimum description length (MDL)*
 - *Bayesian marginal likelihood*

Model-based clustering in R

- `mclust` implements multivariate model-based clustering
- Provides an easy interface to fit several parameterizations
- Model comparison with BIC
- Plotting functionality

```
> library(mclust)
```

[illegible]

version 5.4.6

The full model

- We again have a K -mixture model, this time for multivariate continuous variables \mathbf{y} , whose p.d.f., f , is modeled as

$$f(\mathbf{y}) = \sum_{k=1}^K \pi_k \cdot \text{MVN}(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- Where $\boldsymbol{\mu}_k$ is the class-specific mean vector and
- Where $\boldsymbol{\Sigma}_k$ is the class-specific covariance matrix
- Note: no conditional independence assumption by default
- Would correspond to $\boldsymbol{\Sigma}_k$ being a diagonal matrix for each k

Covariance matrices in mclust

- All the models used in mclust are multivariate normal (Gaussian) (key assumption 1)
- Further "juice" is in structure of the covariance matrices, Σ_k
- Would be tedious to restrict the elements of these matrices directly
- Therefore the mclust people use a trick :

$$\Sigma_g = \lambda_g D_g A_g D_g^T.$$

- This trick is very closely related to factor analysis with orthogonal factors, where D would be loading matrix, and A a diagonal "relative factor variance" matrix, with overall amount of variance lambda

Covariance matrices in `mclust`

"Volume-Shape-Orientation" decomposition:

$$\Sigma_g = \lambda_g D_g A_g D_g^T.$$

- λ_g : volume
 - A_g : shape
 - D_g : orientation
-
- Each of these can be equal (E) or different (V) across classes
 - In addition, A_g and/or D_g can be "identity" (I)

Model-based clustering in R

- Mclust uses an identifier for each possible parametrization :
- **E** for **e**qual, **V** for **v**ariable, **I** for identity matrix:

- **Volume** (size of the clusters in data space):
- **Shape** (circle or ellipse)
- **Orientation** (the angle of the ellipse)



- E.g. an “EEE” model has equal volume, shape and orientation
- A VVV model has variable volume, shape, and orientation
- A VVE model has variable volume and shape but equal orientation

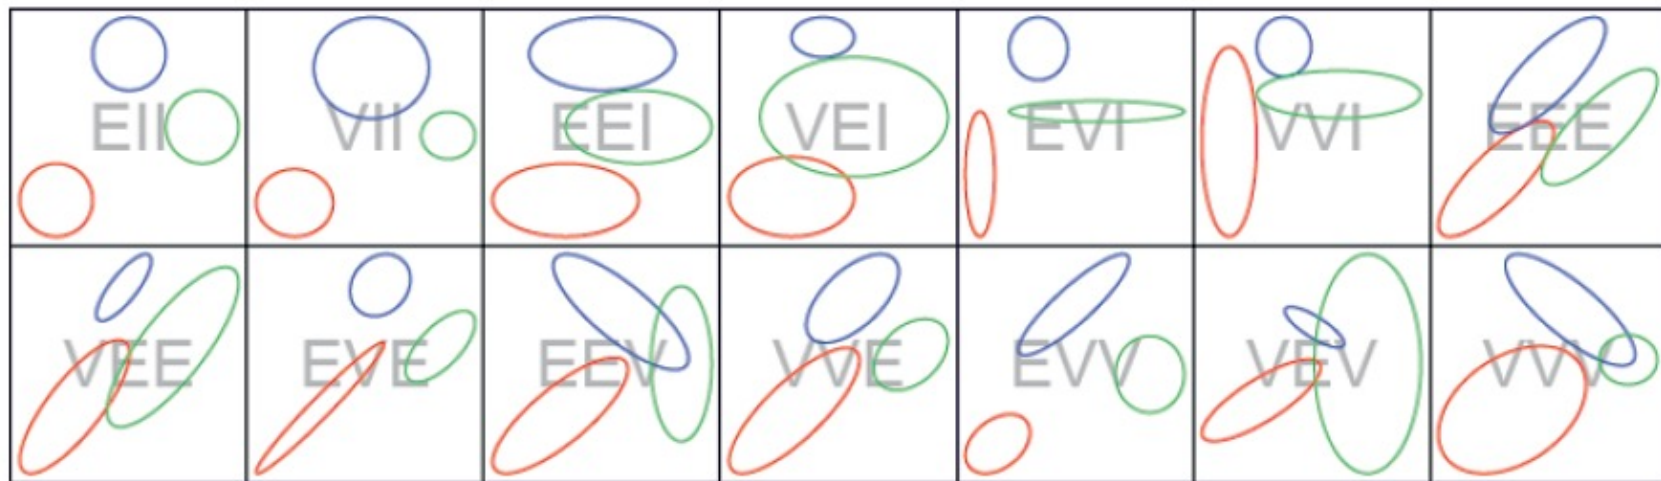


Figure 2.3 Models used in model-based clustering: examples of contours of the bivariate normal component densities for the 14 parameterizations of the covariance matrix used in model-based clustering.

Source: Bouveyron et al. (2021)


Identifier	Model	Distribution	Volume	Shape	Orientation
E V		Univariate Univariate	Equal Variable		
EII VII	λI $\lambda_g I$	Spherical Spherical	Equal Variable	Equal Equal	NA NA
EEI VEI EVI VVI	λA $\lambda_g A$ λA_g $\lambda_g A_g$	Diagonal Diagonal Diagonal Diagonal	Equal Variable Equal Variable	Equal Equal Variable Variable	Axis-aligned Axis-aligned Axis-aligned Axis-aligned
EEE VEE EVE EEV	Σ $\lambda_g D A D^T$ $\lambda D A_g D^T$ $\lambda D_g A D_g^T$	Ellipsoidal Ellipsoidal Ellipsoidal Ellipsoidal	Equal Variable Equal Equal	Equal Equal Variable Equal	Equal Equal Equal Variable
VVE EVV VEV VVV	$\lambda_g D A_g D^T$ $\lambda D_g A_g D_g^T$ $\lambda_g D_g A D_g^T$ Σ_g	Ellipsoidal Ellipsoidal Ellipsoidal Ellipsoidal	Variable Equal Variable Variable	Variable Variable Equal Variable	Equal Variable Variable Variable

Source: Bouveyron et al. (2021)

Model-based clustering in R:

EEE


Equal volume, shape, orientation



vs.

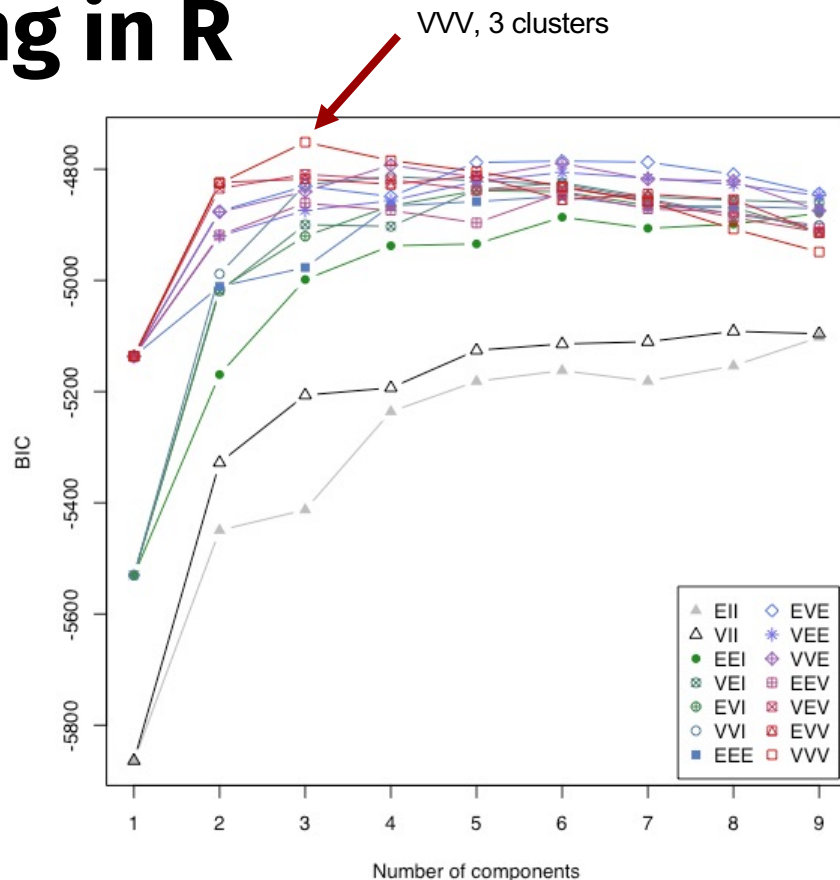
VVV

Variable volume, shape, orientation



Model-based clustering in R

- How `mclust` optimizes hyperparameters:
 - Fit all the models with up to 9 clusters (or more, your choice!)
 - Compute the BIC (or ICL) of each model
 - Choose the model with the best BIC




```
> fit_mc <- Mclust(im_ar, G = 1:10)
fitting ...
|=====| 100%
```

```
> summary(fit_mc)
```

```
-----
Gaussian finite mixture model fitted by EM algorithm
-----
```

```
Mclust VVV (ellipsoidal, varying volume, shape, and orientation)
model with 8 components:
```

log-likelihood	n	df	BIC	ICL
3808542	640000	79	7616028	7530927

```
Clustering table:
```

1	2	3	4	5	6	7	8
151032	48661	155542	34602	82621	49494	41665	76383

GMM in Latent GOLD


- Latent GOLD does not expose the same covariance structure parametrizations as mclust does.
- It implicitly uses a specific kind of parametrization by assuming a diagonal covariance matrix by default, meaning:
 - No correlation between variables within components.
 - Each component has independent variances for each variable.
 - Equivalent to “VVI” in mclust

What if “clusters” are not normally distributed?


What if our “clusters” are not normally distributed?

- Example:
- At this point we could distinguish between:
 - a. “clusters”: the groups we are actually interested in
 - b. “mixture components”: the groups we find when assuming multivariate normal distributions (Gaussians)
- What can we do when clusters \neq components?
- Two ideas discussed here:
 1. Use ICL or entropy directly to select classes
 2. Use component merging


BIC-best classification




BIC-best density



ICL-best classification



ICL-best density



Using entropy (e.g. ICL) to select K

- Provided there is good separation between our clusters,
- Using ICL (entropy) to select number of “classes” gives:
 - The right number of components, and
 - The components more or less correspond to the clusters
- Disadvantage: the model does not fit very well
- (density is not well estimated)
- Can matter for generative models

Merging *components* to get *clusters*

- GMM obviously has trouble with clusters that are not **ellipses**
- Secret weapon: **merging**

Powerful idea:

- Start out with the usual Gaussian mixture solution;
- **merge** “similar” *components* to create non-Gaussian *clusters*.


Note: we’re distinguishing “components” from “clusters” now.

Merging components to get clusters


```
library(mclust)
```

```
output <- clustCombi(data = x)
```


```
plot(output)
```




BIC solution (8 clusters)




Combined solution with 7 clusters




Combined solution with 6 clusters




Combined solution with 5 clusters




Combined solution with 4 clusters




Combined solution with 3 clusters




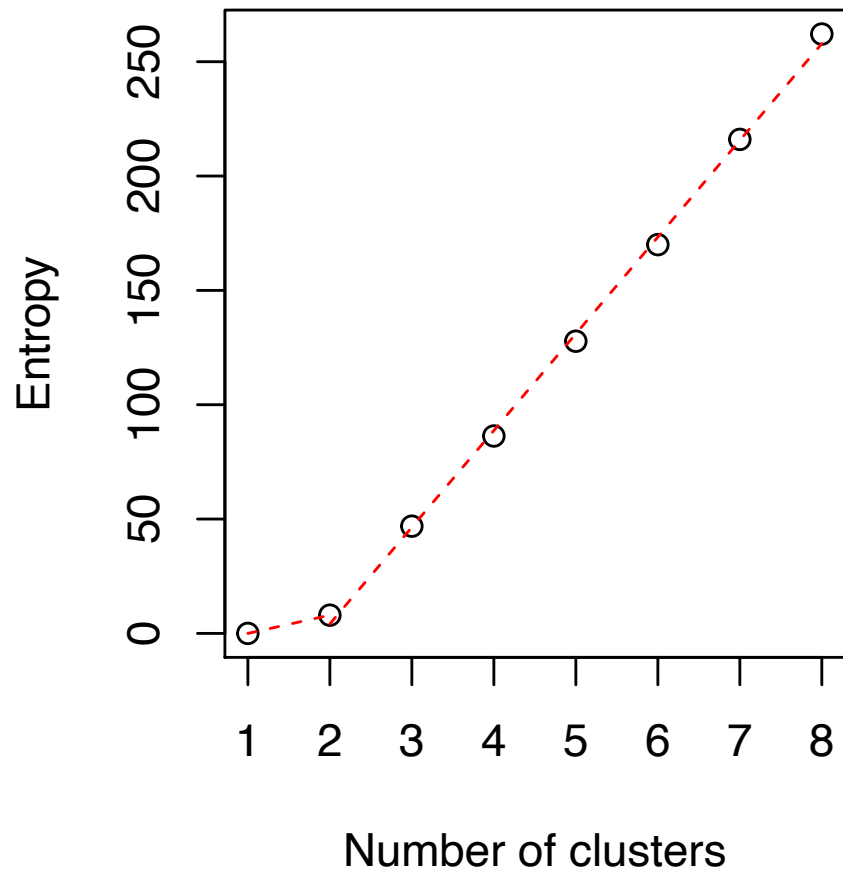
Combined solution with 2 clusters



Combined solution with 1 clusters



Entropy plot



The tidyLPA package

```
library(tidyLPA)
library(dplyr)
```

```
pisaUSA15[1:100, ] %>%
  select(broad_interest, enjoyment, self_efficacy) %>%
  single_imputation() %>%
  estimate_profiles(3)
#> tidyLPA analysis using mclust:
#>
#>   Model Classes AIC      BIC      Entropy prob_min prob_max n_min n_max BLRT_p
#>   1      3      639.57 676.04 0.71      0.67      0.91      0.11 0.60 0.06
```