

Latent class analysis

Latent profile analysis and the EM algorithm

L Boeschoten & DL Oberski

Learning goals

- Explain how the EM algorithm works
- Understand and run R code implementing EM for a two-component mixture of univariate Gaussians
- Explain the key assumption that makes LPA possible





Fig. 1 Peoples' height.

Le Petit Prince



*Mon dessin ne représentait pas un chapeau. Il représentait
un serpent boa qui digérait un éléphant*



*J'ai alors dessiné
l'intérieur du serpent boa, afin que les grandes personnes puissent
comprendre. Elles ont toujours besoin d'explications*



There is a much quoted story about David Hilbert, who one day noticed that a certain student had stopped attending class. When told that the student had decided to drop mathematics to become a poet, Hilbert replied,

“Good — he did not have enough imagination to become a mathematician.”

— Robert Osserman (US mathematician, 1926 – 2011)

Gaussian mixture modeling

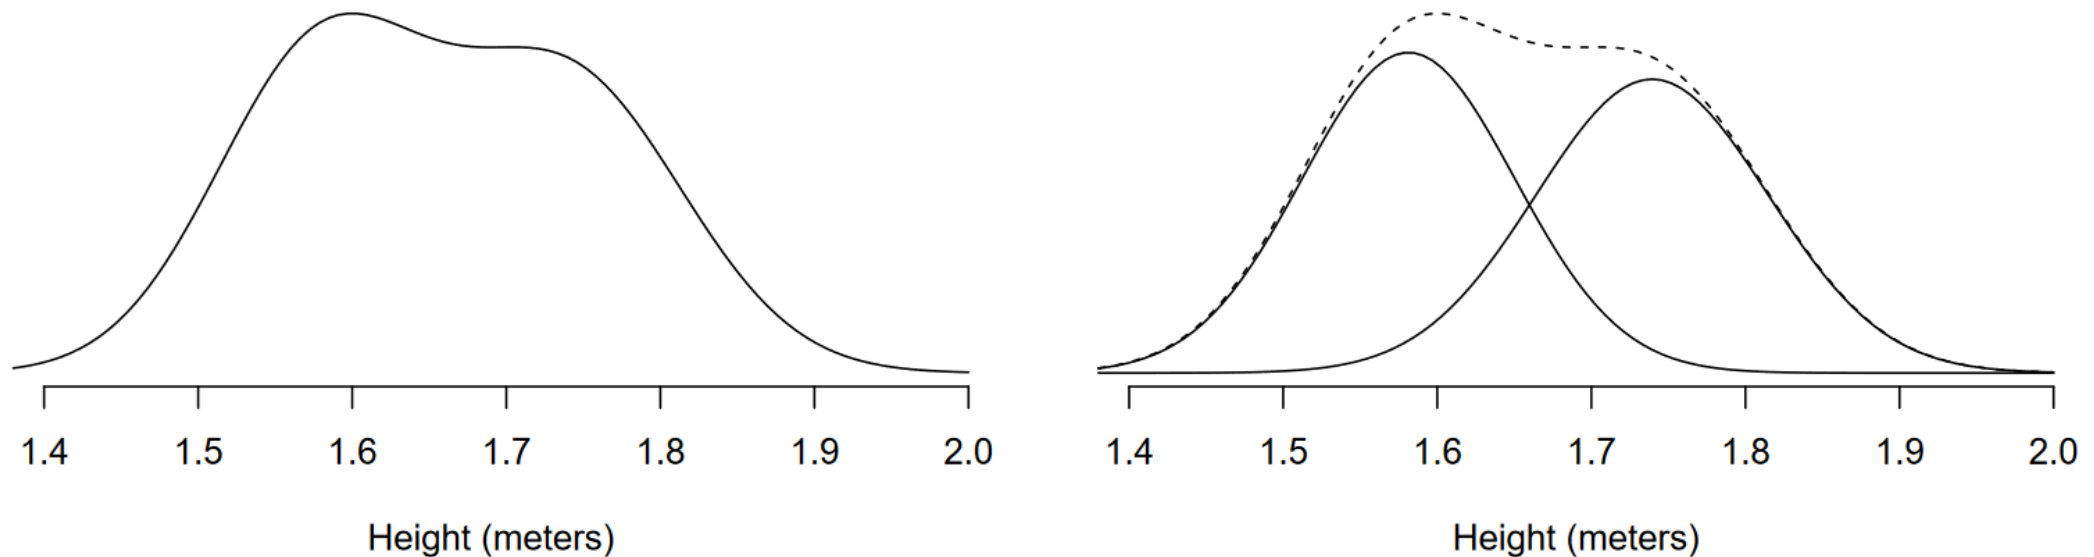


Fig. 1 Peoples' height. Left: observed distribution. Right: men and women separate, with the total shown as a dotted line.

How data were generated

```
n <- 1000
```

```
true_mean_men <- 1.74#m
```

```
true_mean_women <- 1.58#m
```

```
true_sd_men <- 0.08#m
```

```
true_sd_women <- 0.07#m
```

```
true_sex <- rbinom(n, size = 1, prob = 0.5) + 1
```


How data were generated

```
means_mix <- c(true_mean_men,  
               true_mean_women)[true_sex]
```

```
sds_mix <- c(true_sd_men, true_sd_women)[true_sex]
```

```
height <- rnorm(n, means_mix, sds_mix)
```

Gaussian mixture modeling




Fig. 1 Peoples' height. Left: observed distribution. Right: men and women separate, with the total shown as a dotted line.

What mixture modeling does

- Given only the **observed data** (of course!),
- can we recover the values:

```
true_mean_men      <- 1.74#m  
true_mean_women    <- 1.58#m  
true_sd_men        <- 0.08#m  
true_sd_women      <- 0.07#m
```

?

“Maximum likelihood”

- Statistical model = **assumptions** defines a **likelihood**

$$p(\text{data} \mid \text{parameters}) = p(y \mid \theta)$$

- Maximum **likelihood** estimation: find the parameters θ that make it most likely to observe the data we actually observed, y
- The above procedure automatically gives algorithm for computing clusters from data, given the model.

Univariate mixture of Gaussians

Likelihood (*density*) for height data:

$$p(\text{height} \mid \theta) =$$

$$\Pr(\text{man}) \cdot \text{Normal}(\mu_{\text{man}}, \sigma_{\text{man}}) +$$


$$\Pr(\text{woman}) \cdot \text{Normal}(\mu_{\text{woman}}, \sigma_{\text{woman}})$$

Or, more concise notation:

$$p(\text{height} \mid \theta) =$$

$$\pi_1^X \text{Normal}(\mu_1, \sigma_1) +$$

$$(1 - \pi_1^X) \text{Normal}(\mu_2, \sigma_2)$$



Mixture model

Gaussian mixture model **parameters**:

- π_1^X determines the relative cluster sizes
 - Proportion of observations to be expected in each cluster
- μ_1 and μ_2 determine the locations of the clusters
 - Like centroids in K-means clustering
- σ_1 and σ_2 determine the volume of the clusters
 - how large / spread out the are clusters are in data space

Together, these 5 unknown parameters describe our model of how the data is generated.

Plan of attack: direct maximum likelihood

- Write down how we think the data were generated (*imagination step*)
→ **model**
- This involves **parameters** (the means, sds, and class sizes)
- Try out some values for the parameters (*exploration step*)
- Work out what data *would* look like, if our model were right, and the parameters had the values we are currently guessing
- Compare to the real data (*evaluation step*)
- Pick the guessed parameter values that would give data closest to the observed data (*optimization step*)

Why not this plan of attack?

- It gives a general solution to the problem, BUT
- The exploration space is large
- **Difficult** to find a way that always gives the right answer in the end
- Model can be complicated, no ready-made solutions


Estimation: the EM algorithm

- If we knew in advance who is a man and who is a woman, it would have been easy to find the estimates for μ and σ :

$$\hat{\mu}_1 = \frac{\sum_{i=1}^{N_1} \text{height}_i}{N_1}, \quad \hat{\sigma}_1 = \sqrt{\frac{\sum_{i=1}^{N_1} (\text{height}_i - \hat{\mu}_1)^2}{N_1}}$$

(and same for $\hat{\mu}_2$ and $\hat{\sigma}_2$.)


- But we don't know this!
- > **Assignments need to be estimated too.**



Estimation: the EM algorithm

- Solution: Figure out the **posterior probability** of being a man/woman, given the current estimates of the means and sds
- If we know cluster locations and shapes, how likely is it that a 1.7m person is a man or a woman?


$$\pi_{man}^X = \frac{2.20}{2.86} \approx 0.77$$




Estimation: the Expectation-Maximization (EM) algorithm

- Now we have some class *assignments* (probabilities);
- So we can go back to the parameters and update them using our easy rule (M-step)
- Then, we can compute new posterior probabilities (E-step)

Estimation: the EM algorithm



Estimation: the EM algorithm



R code: starting values

```
guess_mean_men <- mean(height) + sep_start  
guess_mean_wom <- mean(height) - sep_start
```

```
guess_sd_men <- sd(height)  
guess_sd_wom <- sd(height)
```

R code: iteration between E and M steps

```
for(it in 1:maxit) {  
  # E-step  
  
  # M-step  
}
```


R code: E-step

```
pman <- dnorm(height, mean = guess_mean_men,  
               sd = guess_sd_men)
```

```
pwom <- dnorm(height, mean = guess_mean_wom,  
               sd = guess_sd_wom)
```

```
post <- pman / (pman + pwom)
```

R code: M-step for means

```
guess_mean_men <-  
  weighted.mean(height, w = post)  
  
guess_mean_wom <-  
  weighted.mean(height, w = 1-post)
```

R code: M-step for sd's

```
guess_sd_men <- sqrt(weighted.mean(  
  (height - guess_mean_men)^2, w = post))
```

```
guess_sd_wom <- sqrt(weighted.mean(  
  (height - guess_mean_wom)^2, w = 1-post))
```

... that's it!

Our simple code works

```
> runit()
```



Iter:	M:	F:	sd(M):	sd(F):

Start	1.86	1.46	0.107	0.107
1	1.74	1.58	0.072	0.066
2	1.74	1.58	0.072	0.066
3	1.74	1.58	0.073	0.066

What just happened?

Went from left to right.

MAGIC???




What just happened?

- We were able to achieve our goal: guess the parameters of the imagined model, using only observed data
- The price we paid was an **assumption**:
Within each group, the distribution of heights is exactly Normal (Gaussian)
- By making this assumption, we could perform our E and M steps


(Mixture) modeling is not based on MAGIC but on ASSUMPTIONS

EM in general

- There is some “missing data” (in our case *class membership*)
- If we had the missing data, finding the parameter estimates would be easy
 - This easy rule is our “maximization” (M) step
- If we had the parameters, then estimating the missing data would be easy
 - This easy rule is our “expectation” (E) step
- Now we just alternate between easy rules – easy!



Model parameter



Source: Bettina Grün

Latent data parameter

EM algorithm

Source: Bettina Grün

Advantages:

- The likelihood is increased in each step → converges
- Relatively easy to implement:
- Different mixture models require only different M-steps.
- Weighted ML estimation of the component specific model is sometimes already available in standard software.

Disadvantages:

- Standard errors have to be determined separately
- Slow convergence.
- **Convergence only to a *local optimum*.**

Convergence to local optimum ☹️

```
> library(flexmix)
> height_fit_fm <- flexmix(height ~ 1, k = 2)
> parameters(height_fit_fm)
```

	Comp.1	Comp.2
coef.(Intercept)	1.6587925	1.6610152
sigma	0.1074166	0.1075208

EM algorithm

- The big disadvantage of EM is that it does not always find the best solution
- Currently, the advice to avoid this problem is to try many different random starting values
- The solution with the best likelihood should be chosen
- And, ideally, should be found several times from different starting points

Multiple random starts solve problem ^^

```
> library(flexmix)
> height_fit_fm <-
  stepFlexmix(height ~ 1, k = 2, nrep = 50)
2 : * * * * * * * * * * * * * * * * * * * * * * * * *
> parameters(height_fit_flexmix)
```


	Comp.1	Comp.2
coef.(Intercept)	1.73240449	1.569787
sigma	0.07717027	0.061939

The classic

Dempster, A. P., Laird, N. M., & Rubin, D. B. (1977). Maximum Likelihood from Incomplete Data via the EM Algorithm. *Journal of the Royal Statistical Society. Series B (Methodological)*, 39(1), 1–38.

The EM Algorithm and Extensions

Second Edition



Geoffrey J. McLachlan
Thriyambakam Krishnan

WWW.
WILEY-SON